

# High Performance Parallel Computing

## Assignment 2

António Maschio ; Dimitrios Anastasiou ; Georgios Sevastakis

February 2024

### 1 Objective

The objective is to optimise a given code on molecular dynamics by vectorising the code and subsequently, by using the OpenMP API (Application programming interface). A profiler was implemented to investigate the performance and pragmas were added to force vectorization the code.

### 2 Introduction

In this assignment, we were given a code simulating a system of  $N$  water molecules (their interactions) carried through with the leap-frog integrator. In this  $N$ -body simulation, the time-boost four potentials were included; For atoms belonging to the same molecule, there were the bond potential and the angle potential simulating the oscillating fashion of the bond and the angle, accordingly, while for the rest of the atoms, the Lennard-Jones (LJ) and the Coulomb potentials were taken into account.

As far as the LJ potential is concerned, the  $r^{-12}$  term models the strong repulsive forces indicated by the Pauli exclusion principle, while the  $r^{-6}$  term models the Van der Waals force. Regarding the Coulomb potential, it is interesting to note that in many MD models the charges can be less than 1. This is due to the fact that we are dealing with bounded particles that cannot exist as free particles (like quarks) or that we are dealing with quasiparticles, meaning particles that are not real but they behave like if they were.

### 3 Task 1: Create a Struct-of-Arrays version of the program (amenable to vectorization)

In order to transition to Struct-of-Arrays the loops were modified accordingly. For example, the below code snippet shows the four for loops of the UpdateNonBondedForces function that they now iterate through the position, velocity and force vectors instead of the different molecules. The way the data are accessed can also be seen in the last line of the snippet. The complete code can be found in [Appendix A](#).

```
1 for (auto& atom1 : sys.molecules.atoms){  
2     for (auto& atom2 : sys.molecules.atoms){  
3         for (long unsigned int i = 0; i < sys.molecules.no_mol; i++){  
4             for (long unsigned int j = i+1; j < sys.molecules.no_mol; j++){  
5                 Vec3 dp = atom1.p[i]-atom2.p[j];
```

Modified UpdateNonBondedForces function in the SoA format.

## 4 Task 2: Investigate the performance of the code with a profiler

### 4.1 Using 4 molecules, which functions contribute most to the runtime?

When running the simulation for 100000 steps and 4 molecules, the following profiling was obtained using the gprof profiler:

% Time	Name
83.33	UpdateNonBondedForces
11.11	UpdateAngleForces
5.56	UpdateBondForces
0.00	Evolve
0.00	Vector allocator
0.00	Write
0.00	Make Water

Table 1: Function Time Percentages for the case of 4 molecules with steps = 100000.

From this table, we can deduce that the functions that contribute the most to the runtime of this simulation are the UpdateNonBondedForces and UpdateAngleForces.

### 4.2 How does it change if modeling 2, 16, and 128 molecules? Explain why the percentages change.

The results obtained are summarized in the following table:

Name	2 mol.	4 mol.	16 mol.	128 mol.
UpdateNonBondedForces	100%	63.64%	94.02%	99.26%
UpdateAngleForces	-	36.36%	3.42%	0.45%
UpdateBondForces	-	-	1.71%	0.19%
Evolve	-	-	0.85%	0.09%

Table 2: Function Time percentages for 4 different simulations with steps = 100000.

We need to look at how the code works to understand why the percentages change as seen at the table above. First of all, let's have a look at the for-loops of the individual functions, since they contribute the most to the complexity of the code:

Name	Number of iterations
UpdateNonBondedForces	$A^2 \cdot \frac{M(M-1)}{2}$
UpdateAngleForces	$M$
UpdateBondForces	$B \cdot M$
Evolve	$A \cdot M$

Table 3: Number of iterations of each individual function, where  $M$ ,  $A$  and  $B$  stand for the number of molecules, atoms and bonds respectively.

In the code, the number of atoms  $A$  and that of the bonds  $B$  remain constant and are equal to 3 and 2 respectively, meaning that all functions but the UpdateNonBondedForces run  $M$  times, while the UpdateNonBondedForces run  $M^2$  times. This justifies the prevalence of the UpdateNonBondedForces function for larger  $M$  values.

In the  $M = 2$  and  $M = 4$  cases, the number of iterations is not the most important factor contributing

to the complexity. In the former one, UpdateNonBondedForces completely dominates due to the  $A^2$  term (100% of the elapsed time), while in the latter, it prevails due to the combination of  $M$  and  $A$  ( $\approx 60\%$  of the elapsed time). For  $M = 4$ , we also notice that 1/3 of the elapsed time is consumed by the UpdateAngleForces. This is due to the main body of this particular function; it includes many function calls and declarations and that is also the reason it preserves the second place for greater  $M$ -values. As far as the UpdateBondForces and Evolve functions are concerned, they share the same number of iterations, but the Evolve is simpler than UpdateBondForces, resulting in the former coming last.

### 4.3 Which function is most important to get good performance

The most important function to be optimized is the UpdateNonBondedForces as stated in the previous subsection for the aforementioned reasons.

### 4.4 How does the vectorized (SoA) version perform compared to the original sequential (AoS) version running with 2 molecules and with 128 molecules. If there is a difference in the relative performance, please discuss why this could be the case.

Name	2 mol. Seq	2 mol. vec	128 mol. seq	128 mol. vec
UpdateNonBondedForces	0.03 s	0.01	162.60 s	173.28 s
UpdateAngleForces	-	0.01	0.57 s	0.72 s
UpdateBondForces	0.02 s	0.03	0.58 s	0.38 s
Evolve	0.01 s	-	0.29 s	0.15 s
<b>Total</b>	0.05646 s	0.05794	164.4 s	174.9 s

Table 4: Time in seconds for each simulation. steps = 100000

From the table above, we notice that the vectorized version performs a bit worse than the sequential one and Memory access patterns is to blame.

The sequential code was making good use of the cache since with one memory call more data would be stored in cache, which would subsequently be called right away. On the other hand, in the vectorised version the SoA format resulted in many cache-misses and therefore, worse performance.

## 5 Adding OpenMP SIMD pragmas to the code

The pragmas were placed according to the results of the code profiler. The most time consuming functions of our code are the ones updating the forces along with the Evolve function. We chose to place the pragmas on those functions and on the loops that iterate through each molecule since the operation done inside these loops can be vectorised with simd (i.e. updating multiple data with one instruction). The clauses of the pragmas were the reduction used on the accumulated forces and collapse of the two for loops in the UpdateNonBondedForces function.

# Appendices

## A C++ code

```

1 #include <iostream>
2 #include <iomanip>
3 #include <fstream>
4 #include <vector>
5 #include <cassert>
6 #include <math.h>
7 #include <chrono>
8 #include <list>
9
10 const double deg2rad = acos(-1)/180.0; // pi/180 for changing degs to radians
11 double accumulated_forces_bond = 0.; // Checksum: accumulated size of forces
12 double accumulated_forces_angle = 0.; // Checksum: accumulated size of forces
13 double accumulated_forces_non_bond = 0.; // Checksum: accumulated size of forces
14
15 class Vec3 {
16 public:
17     double x, y, z;
18     // initialization of vector
19     Vec3(double x, double y, double z): x(x), y(y), z(z) {}
20     // size of vector
21     double mag() const{
22         return sqrt(x*x+y*y+z*z);
23     }
24     Vec3 operator-(const Vec3& other) const{
25         return {x - other.x, y - other.y, z - other.z};
26     }
27     Vec3 operator+(const Vec3& other) const{
28         return {x + other.x, y + other.y, z + other.z};
29     }
30     Vec3 operator*(double scalar) const{
31         return {scalar*x, scalar*y, scalar*z};
32     }
33     Vec3 operator/(double scalar) const{
34         return {x/scalar, y/scalar, z/scalar};
35     }
36     Vec3& operator+=(const Vec3& other){
37         x += other.x; y += other.y; z += other.z;
38         return *this;
39     }
40     Vec3& operator-=(const Vec3& other){
41         x -= other.x; y -= other.y; z -= other.z;
42         return *this;
43     }
44     Vec3& operator*=(double scalar){
45         x *= scalar; y *= scalar; z *= scalar;
46         return *this;
47     }
48     Vec3& operator/=(double scalar){
49         x /= scalar; y /= scalar; z /= scalar;
50         return *this;
51     }
52 };
53 Vec3 operator*(double scalar, const Vec3& y){
54     return y*scalar;
55 }
56 Vec3 cross(const Vec3& a, const Vec3& b){
57     return { a.y*b.z-a.z*b.y,

```

```

58         a.z*b.x-a.x*b.z,
59         a.x*b.y-a.y*b.x };
60 }
61 double dot(const Vec3& a, const Vec3& b){
62     return a.x * b.x + a.y * b.y + a.z * b.z;
63 }
64
65 /* a class for the bond between two atoms  $U = 0.5k(r_{12}-L_0)^2$  */
66 class Bond {
67 public:
68     double K;        // force constant
69     double L0;       // relaxed length
70     int a1, a2;      // the indexes of the atoms at either end
71 };
72
73 /* a class for the angle between three atoms  $U=0.5K(\phi_{123}-\phi_0)^2$  */
74 class Angle {
75 public:
76     double K;
77     double Phi0;
78     int a1, a2, a3; // the indexes of the three atoms, with a2 being the centre atom
79 };
80
81 // =====
82 // Two newclasses arranging Atoms in a Structure-of-Array data structure
83 // =====
84
85 /* atom class, represent N instances of identical atoms */
86 class Atoms {
87 public:
88     // The mass of the atom in (U)
89     double mass;
90     double ep;           // epsilon for LJ potential
91     double sigma;        // Sigma, somehow the size of the atom
92     double charge;       // charge of the atom (partial charge)
93     std::string name;    // Name of the atom
94     // the position in (nm), velocity (nm/ps) and forces (k_BT/nm) of the atom
95     std::vector<Vec3> p,v,f;
96     // constructor, takes parameters and allocates p, v and f properly to have
97     // N_identical elements
98     Atoms(double mass, double ep, double sigma, double charge, std::string name,
99           size_t N_identical)
100         : mass{mass}, ep{ep}, sigma{sigma}, charge{charge}, name{name},
101           p{N_identical, {0,0,0}}, v{N_identical, {0,0,0}}, f{N_identical, {0,0,0}}
102     {}
103 };
104
105 /* molecule class for no_mol identical molecules */
106 class Molecules {
107 public:
108     std::vector<Atoms> atoms;           // list of atoms in the N identical molecule
109     std::vector<Bond> bonds;            // the bond potentials, eg for water the left
110     // and right bonds
111     std::vector<Angle> angles;          // the angle potentials, for water just the
112     // single one, but keep it a list for generality
113     int no_mol;
114 };
115
116 // =====
117
118 /* system class */
119 class System {

```

```

116 public:
117     Molecules molecules;           // all the molecules in the system
118     double time = 0;               // current simulation time
119 };
120
121 class Sim_Configuration {
122 public:
123     int steps = 10000;             // number of steps
124     int no_mol = 4;               // number of molecules
125     double dt = 0.0005;           // integrator time step
126     int data_period = 100;        // how often to save coordinate to trajectory
127     std::string filename = "trajectory.txt"; // name of the output file with
    trajectory
128     // system box size. for this code these values are only used for vmd, but in
    general md codes, period boundary conditions exist
129
130     // simulation configurations: number of step, number of the molecules in the
    system,
131     // IO frequency, time step and file name
132     Sim_Configuration(std::vector<std::string> argument){
133         for (long unsigned int i = 1; i<argument.size() ; i += 2){
134             std::string arg = argument.at(i);
135             if(arg=="-h"){ // Write help
136                 std::cout << "MD -steps <number of steps> -no_mol <number of
    molecules>"
137
138                 << " -fwrite <io frequency> -dt <size of timestep> -ofile <
    filename> \n";
139                 exit(0);
140                 break;
141             } else if(arg=="-steps"){
142                 steps = std::stoi(argument[i+1]);
143             } else if(arg=="-no_mol"){
144                 no_mol = std::stoi(argument[i+1]);
145             } else if(arg=="-fwrite"){
146                 data_period = std::stoi(argument[i+1]);
147             } else if(arg=="-dt"){
148                 dt = std::stof(argument[i+1]);
149             } else if(arg=="-ofile"){
150                 filename = argument[i+1];
151             } else{
152                 std::cout << "---> error: the argument type is not recognized \n";
153             }
154         }
155
156         dt /= 1.57350; /// convert to ps based on having energy in k_BT, and length
    in nm
157     }
158 };
159 // Given a bond, updates the force on all atoms correspondingly
160
161 ////////////// *****UPDATED***** //////////////////
162
163 void UpdateBondForces(System& sys){
164     // Loops over the (2 for water) bond constraints
165     for (Bond& bond : sys.molecules.bonds){
166         auto& atom1=sys.molecules.atoms[bond.a1];
167         auto& atom2=sys.molecules.atoms[bond.a2];
168         #pragma omp simd reduction(+:accumulated_forces_bond)
169         for(int i=0 ; i < sys.molecules.no_mol ; i++ ){
170             Vec3 dp = atom1.p[i]-atom2.p[i];
171             Vec3 f = -bond.K*(1-bond.L0/dp.mag())*dp;

```

```

172         atom1.f[i] += f;
173         atom2.f[i] -= f;
174         accumulated_forces_bond += f.mag();
175     }
176 }
177 }
178
179 // Iterates over all bonds in molecules (for water only 2: the left and right)
180 // And updates forces on atoms correspondingly
181
182 ////////////// *****UPDATED***** //////////////
183
184 void UpdateAngleForces(System& sys){
185     Angle& angle = sys.molecules.angles[0];
186     auto& atom1=sys.molecules.atoms[angle.a1];
187     auto& atom2=sys.molecules.atoms[angle.a2];
188     auto& atom3=sys.molecules.atoms[angle.a3];
189     #pragma omp simd reduction(+:accumulated_forces_angle)
190     for(int i=0 ; i < sys.molecules.no_mol ; i++ ){
191         //==== angle forces (H--O---H bonds) U_angle = 0.5*k_a(phi-phi_0)^2
192         //f_H1 = K(phi-ph0)/|H10|*Ta
193         // f_H2 = K(phi-ph0)/|H20|*Tc
194         // f_0 = -f1 - f2
195         // Ta = norm(H10 x (H10 x H20))
196         // Tc = norm(H20 x (H20 x H10))
197         //=====
198
199         Vec3 d21 = atom2.p[i]-atom1.p[i];
200         Vec3 d23 = atom2.p[i]-atom3.p[i];
201
202         // phi = d21 dot d23 / |d21| |d23|
203         double norm_d21 = d21.mag();
204         double norm_d23 = d23.mag();
205         double phi = acos(dot(d21, d23) / (norm_d21*norm_d23));
206
207         // d21 cross (d21 cross d23)
208         Vec3 c21_23 = cross(d21, d23);
209         Vec3 Ta = cross(d21, c21_23);
210         Ta /= Ta.mag();
211
212         // d23 cross (d23 cross d21) = - d23 cross (d21 cross d23) = c21_23 cross
213         d23
214         Vec3 Tc = cross(c21_23, d23);
215         Tc /= Tc.mag();
216
217         Vec3 f1 = Ta*(angle.K*(phi-angle.Phi0)/norm_d21);
218         Vec3 f3 = Tc*(angle.K*(phi-angle.Phi0)/norm_d23);
219
220         atom1.f[i] += f1;
221         atom2.f[i] -= f1+f3;
222         atom3.f[i] += f3;
223
224         accumulated_forces_angle += f1.mag() + f3.mag();
225     }
226 }
227
228 // Iterates over all atoms in both molecules
229 // And updates forces on atoms correspondingly
230
231 void UpdateNonBondedForces(System& sys){
232     /* nonbonded forces: only a force between atoms in different molecules
233        The total non-bonded forces come from Lennard Jones (LJ) and coulomb interactions

```

```

233     U = ep[(sigma/r)^12-(sigma/r)^6] + C*q1*q2/r */
234     for (auto& atom1 : sys.molecules.atoms){
235         for (auto& atom2 : sys.molecules.atoms){// iterate over all pairs of atoms,
similar as well as dissimilar
236             #pragma omp simd reduction(+:accumulated_forces_non_bond) collapse(2)
237             for (long unsigned int i = 0; i < sys.molecules.no_mol; i++){
238                 for (long unsigned int j = i+1; j < sys.molecules.no_mol; j++){
239                     Vec3 dp = atom1.p[i]-atom2.p[j];
240
241                     double r = dp.mag();
242                     double r2 = r*r;
243                     double ep = sqrt(atom1.ep*atom2.ep); // ep = sqrt(ep1*ep2)
244                     double sigma = 0.5*(atom1.sigma+atom2.sigma); // sigma = (sigma1
+sigma2)/2
245                     double q1 = atom1.charge;
246                     double q2 = atom2.charge;
247
248                     double sir = sigma*sigma/r2; // crossection**2 times inverse
squared distance
249                     double KC = 80*0.7; // Coulomb prefactor
250                     Vec3 f = ep*(12*pow(sir,6)-6*pow(sir,3))*sir*dp + KC*q1*q2/(r*r2)
*dp; // LJ + Coulomb forces
251                     atom1.f[i] += f;
252                     atom2.f[j] -= f;
253
254                     accumulated_forces_non_bond += f.mag();
255                 }
256             }
257         }
258     }
259 }
260
261 // integrating the system for one time step using Leapfrog symplectic integration
262
263 ////////////// *****UPDATED***** //////////////
264
265 void Evolve(System &sys, Sim_Configuration &sc){
266
267     // Kick velocities and zero forces for next update
268     // Drift positions: Loop over molecules and atoms inside the molecules
269     for (auto& atom : sys.molecules.atoms){
270         #pragma omp simd
271         for(int i=0 ; i < sys.molecules.no_mol ; i++ ){
272             atom.v[i] += sc.dt/atom.mass*atom.f[i]; // Update the velocities
273             atom.f[i] = {0,0,0}; // set the forces zero to prepare
for next potential calculation
274             atom.p[i] += sc.dt* atom.v[i]; // update position
275         }
276     }
277
278     // Update the forces on each particle based on the particles positions
279     // Calculate the intermolecular forces in all molecules
280     UpdateBondForces(sys);
281     UpdateAngleForces(sys);
282     // Calculate the intramolecular LJ and Coulomb potential forces between all
molecules
283     UpdateNonBondedForces(sys);
284
285     sys.time += sc.dt; // update time
286 }
287
288 // Setup one water molecule

```



```

289
290 ////////////// *****UPDATED***** //////////////
291
292 System MakeWater(int N_molecules){
293     //=====
294     // creating water molecules at position X0,Y0,Z0. 3 atoms
295     //              H---O---H
296     // The angle is 104.45 degrees and bond length is 0.09584 nm
297     //=====
298     // mass units of dalton
299     // initial velocity and force is set to zero for all the atoms by the constructor
300     const double L0 = 0.09584;
301     const double angle = 104.45*deg2rad;
302
303     //      mass      ep      sigma charge name
304     Atoms Oatom(16, 0.65,      0.31, -0.82, "O",N_molecules); // Oxygen atoms
305     Atoms Hatom1( 1, 0.18828, 0.238, 0.41, "H",N_molecules); // Hydrogen atoms
306     Atoms Hatom2( 1, 0.18828, 0.238, 0.41, "H",N_molecules);
307
308     // bonds between first H-O and second H-O respectively
309     std::vector<Bond> waterbonds = {
310         { .K = 20000, .L0 = L0, .a1 = 0, .a2 = 1},
311         { .K = 20000, .L0 = L0, .a1 = 0, .a2 = 2}
312     };
313
314     // angle between H-O-H
315     std::vector<Angle> waterangle = {
316         { .K = 1000, .Phi0 = angle, .a1 = 1, .a2 = 0, .a3 = 2 }
317     };
318
319     System sys;
320     #pragma omp simd
321     for (int i = 0; i < N_molecules; i++){
322         Vec3 P0{i * 0.2, i * 0.2, 0};
323         Oatom.p[i] = {P0.x, P0.y, P0.z};
324         Hatom1.p[i] = {P0.x+L0*sin(angle/2), P0.y+L0*cos(angle/2), P0.z};
325         Hatom2.p[i] = {P0.x-L0*sin(angle/2), P0.y+L0*cos(angle/2), P0.z};
326     }
327
328     sys.molecules.atoms.push_back(Oatom);
329     sys.molecules.atoms.push_back(Hatom1);
330     sys.molecules.atoms.push_back(Hatom2);
331     sys.molecules.bonds= waterbonds;
332     sys.molecules.angles= waterangle;
333     sys.molecules.no_mol=N_molecules;
334
335     // Store atoms, bonds and angles in Water class and return
336     return sys;
337 }
338
339 // Write the system configurations in the trajectory file.
340 void WriteOutput(System& sys, std::ofstream& file){
341     // Loop over all atoms in model one molecule at a time and write out position
342     for (auto& atom : sys.molecules.atoms){
343         #pragma omp simd
344         for(int i = 0 ; i < sys.molecules.no_mol ; i++){
345             file << sys.time << " " << atom.name << " "
346                 << atom.p[i].x << " "
347                 << atom.p[i].y << " "
348                 << atom.p[i].z << '\n';
349         }
350     }

```

```

351 }
352
353 //
=====
354 //===== Main function
=====
355 //
=====

356 int main(int argc, char* argv[]){
357     Sim_Configuration sc({argv, argv+argc}); // Load the system configuration from
command line data
358
359     System sys = MakeWater(sc.no_mol); // this will create a system containing sc.
no_mol water molecules
360     std::ofstream file(sc.filename); // open file
361
362     WriteOutput(sys, file); // writing the initial configuration in the trajectory
file
363
364     auto tstart = std::chrono::high_resolution_clock::now(); // start time (nano-
seconds)
365
366     // Molecular dynamics simulation
367     for (int step = 0; step < sc.steps ; step++){
368
369         Evolve(sys, sc); // evolving the system by one step
370         if (step % sc.data_period == 0){
371             //writing the configuration in the trajectory file
372             WriteOutput(sys, file);
373         }
374     }
375
376     auto tend = std::chrono::high_resolution_clock::now(); // end time (nano-seconds)
377
378     std::cout << "Elapsed time:" << std::setw(9) << std::setprecision(4)
<< (tend - tstart).count()*1e-9 << "\n";
379     std::cout << "Accumulated forces Bonds : " << std::setw(9) << std::
setprecision(5)
380         << accumulated_forces_bond << "\n";
381     std::cout << "Accumulated forces Angles : " << std::setw(9) << std::
setprecision(5)
382         << accumulated_forces_angle << "\n";
383     std::cout << "Accumulated forces Non-bond: " << std::setw(9) << std::
setprecision(5)
384         << accumulated_forces_non_bond << "\n";
385 }
386

```

Complete code.